

Stream Ciphers

January 25, 2022



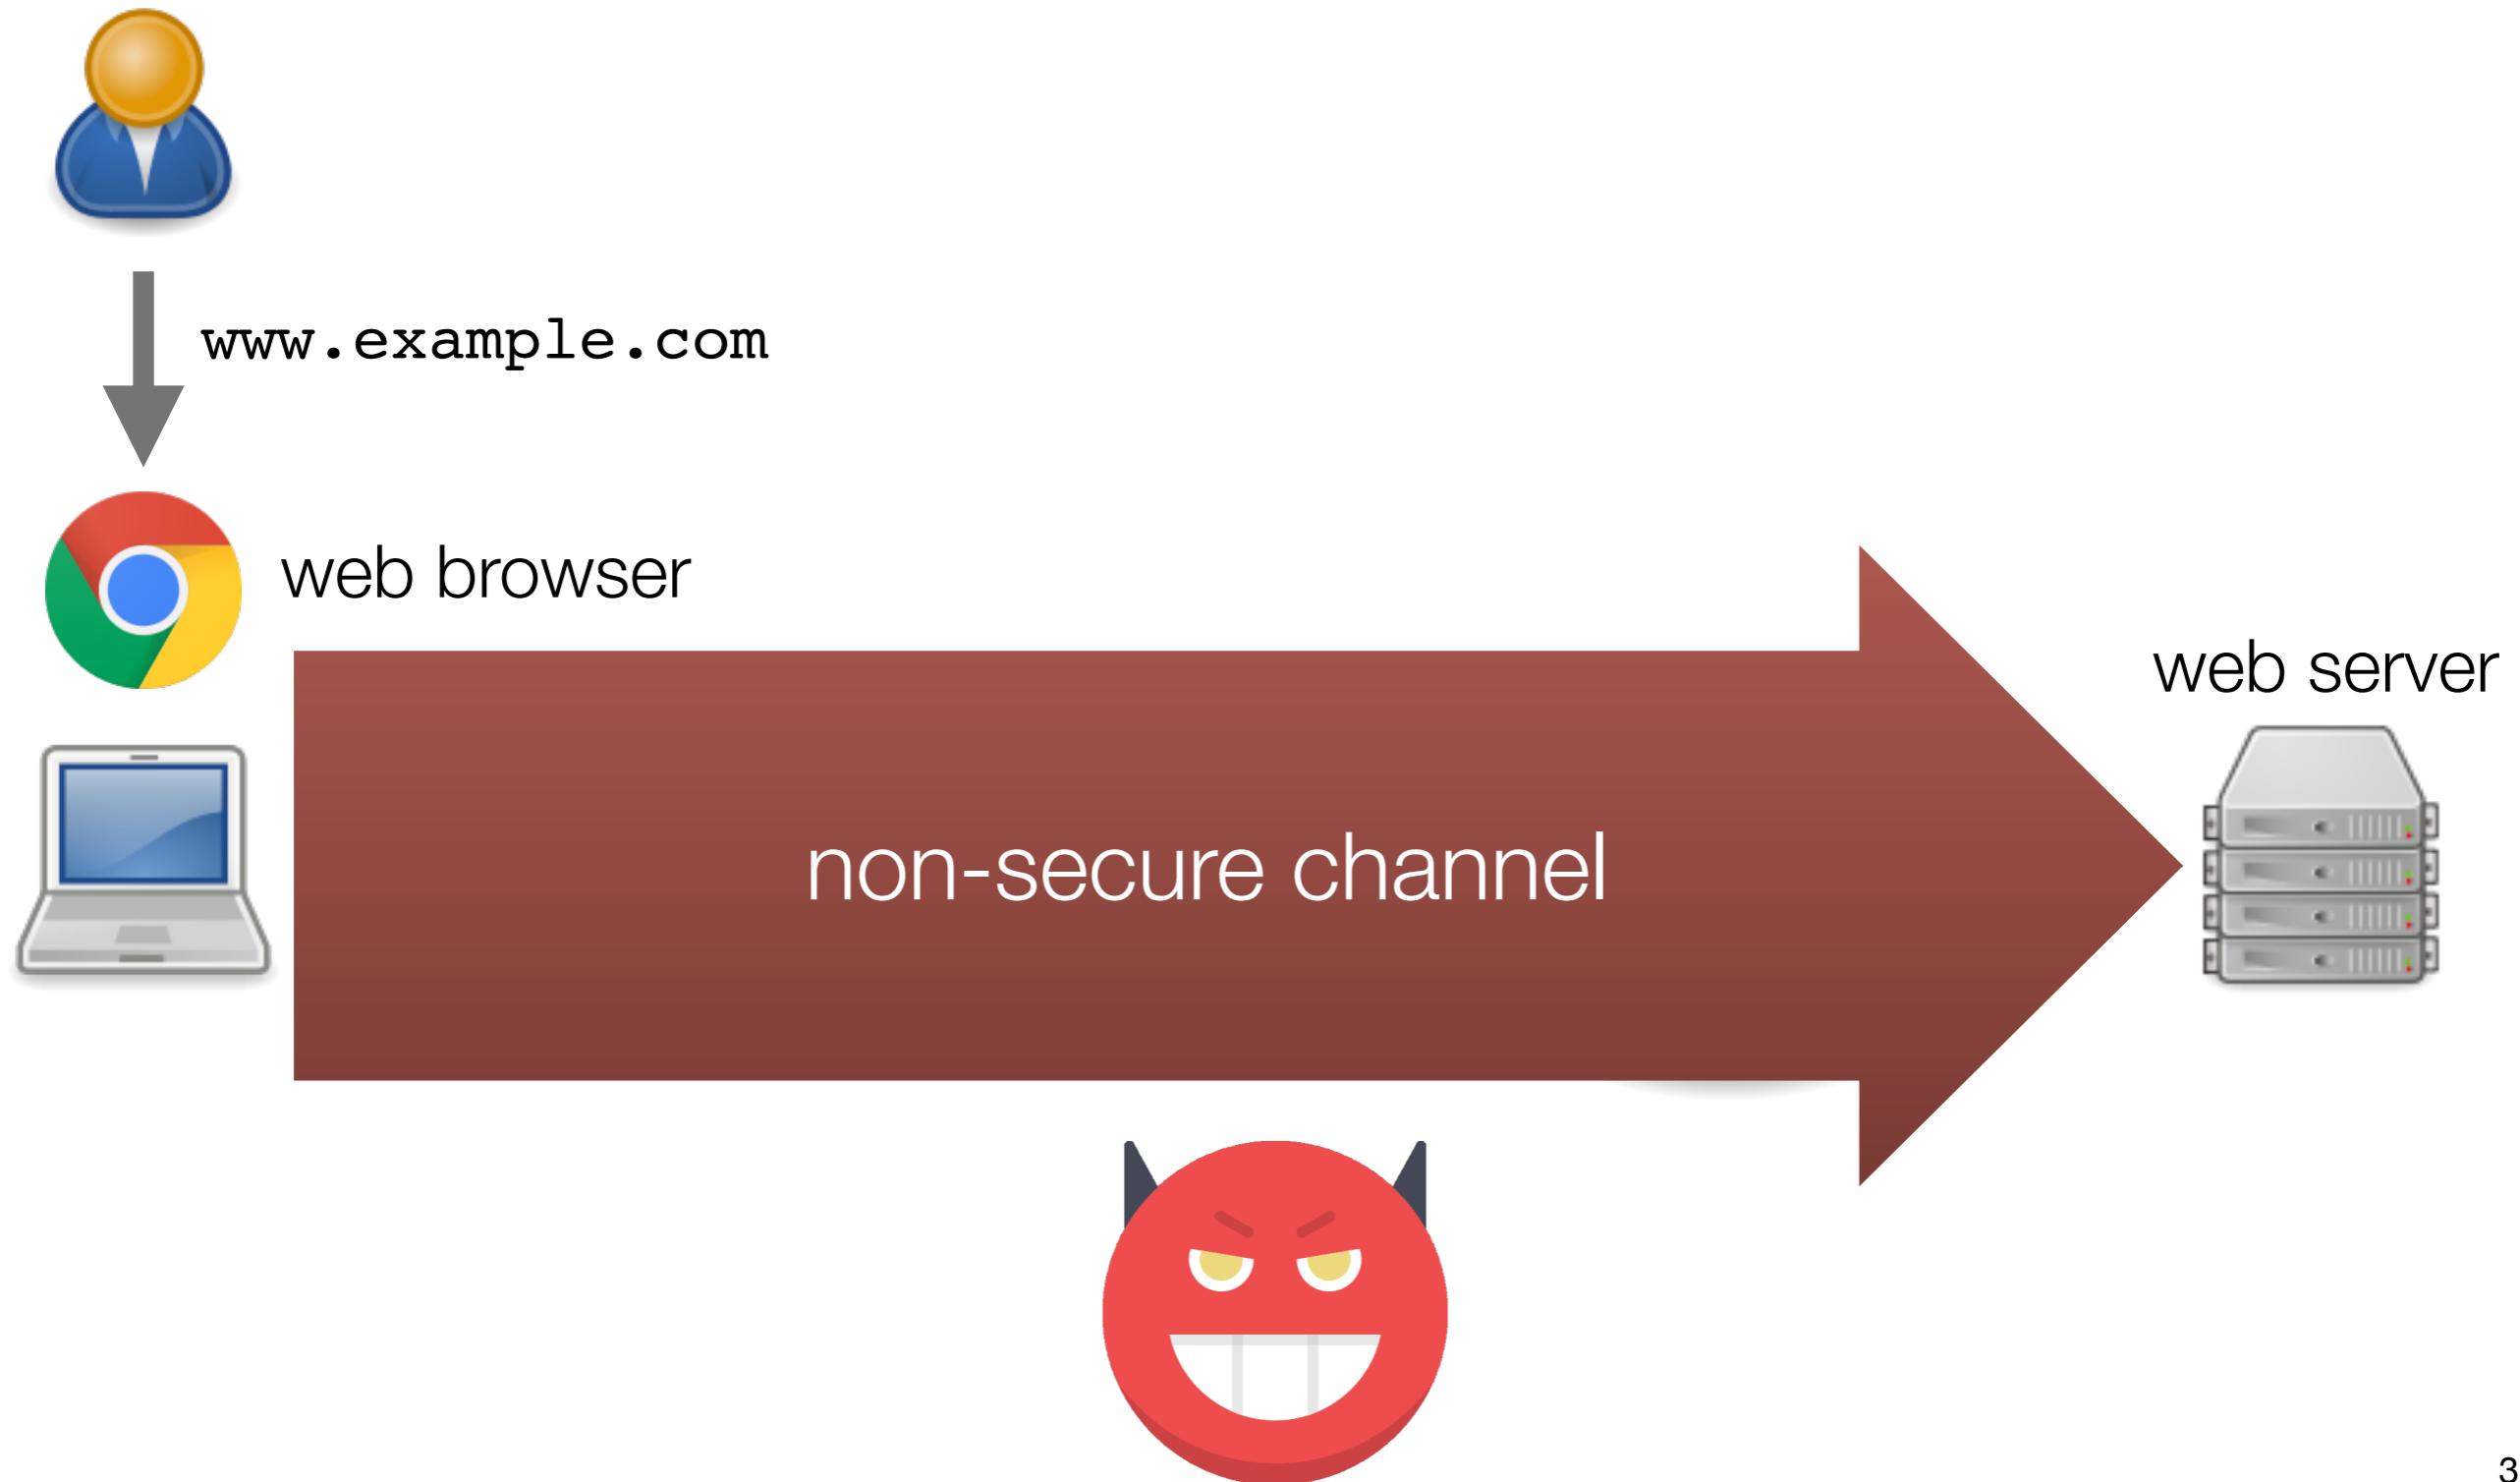
Today

1. perfect security (*but don't get your hopes up, it's not practical*)
2. practical security:
commonly used **stream ciphers** (RC4, Salsa20 / ChaCha20)

Thank you for filling out the background survey online!

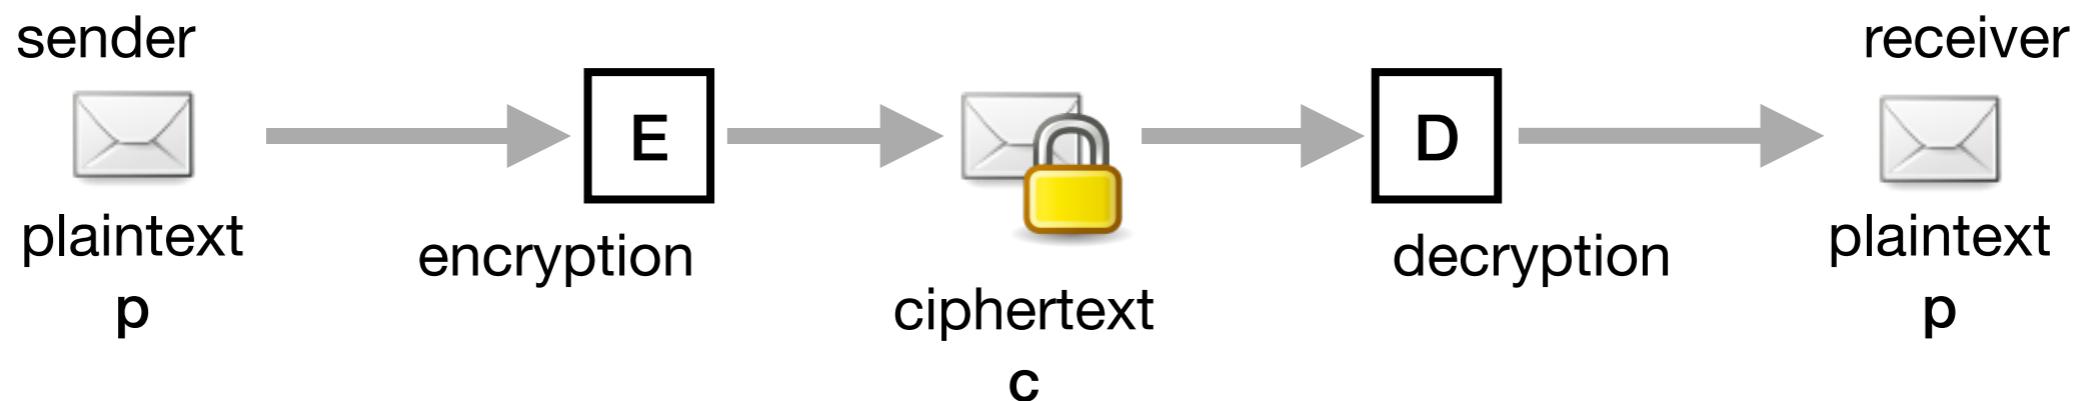
Feedback: <https://forms.gle/JGbNCmCsU69iWaTv8>

Motivating Example: Web Security



Reminder: Encryption

- Model
 - plaintext **p**
 - encryption algorithm (cipher) **E**
 - ciphertext **c**
 - decryption algorithm (cipher) **D**



Kerckhoffs's Principle

- Design requirement for ciphers (1883):
“A cryptographic system should be secure even if all of its details, except for the key, are publicly known.”
- Rejection of security by obscurity for cryptography

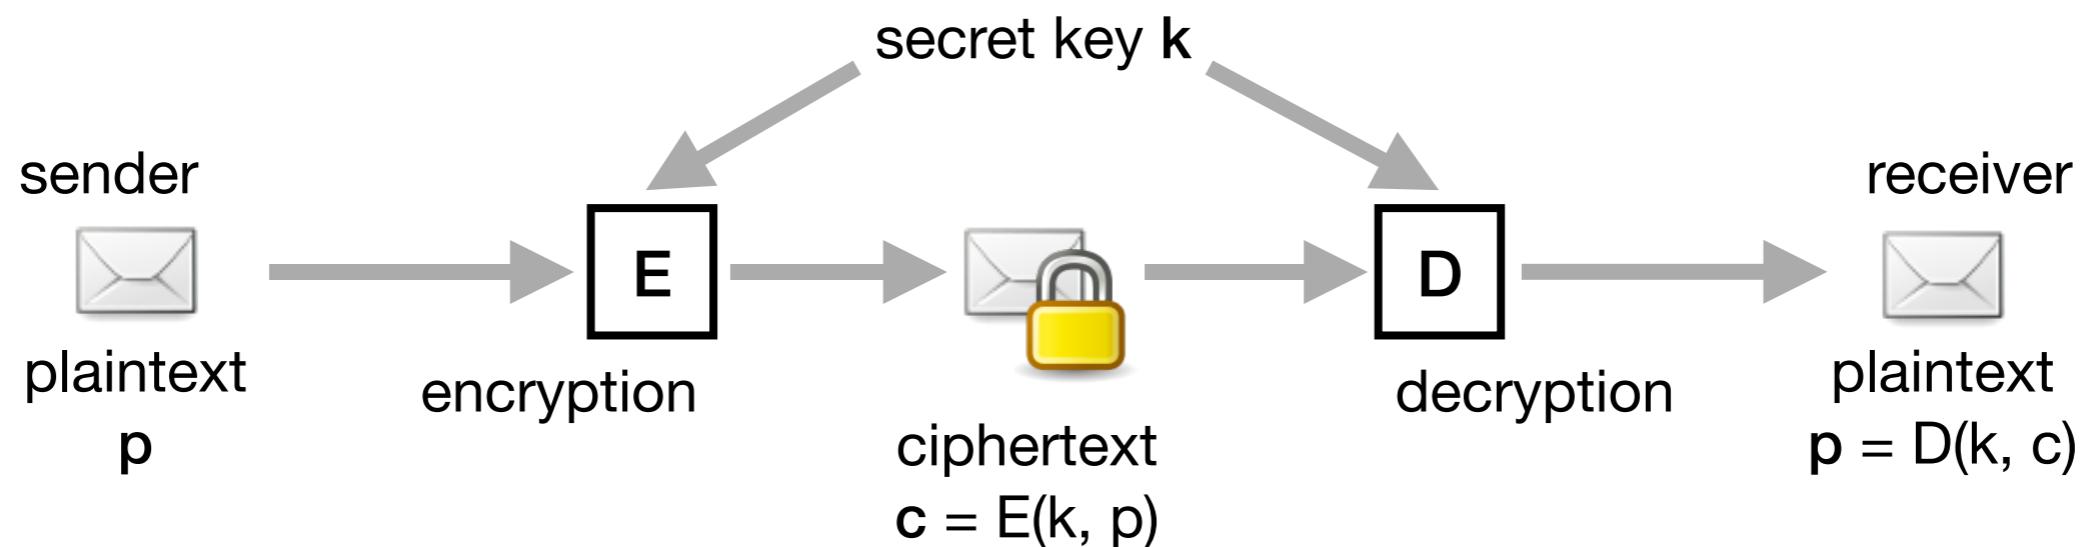


Auguste Kerckhoff

Reminder:

Encryption

- Model
 - plaintext **p**
 - secret key **k**
 - encryption algorithm (cipher) **E**
 - ciphertext **c** = $E(k, p)$
 - decryption algorithm (cipher) **D**



Example Sources of True Randomness

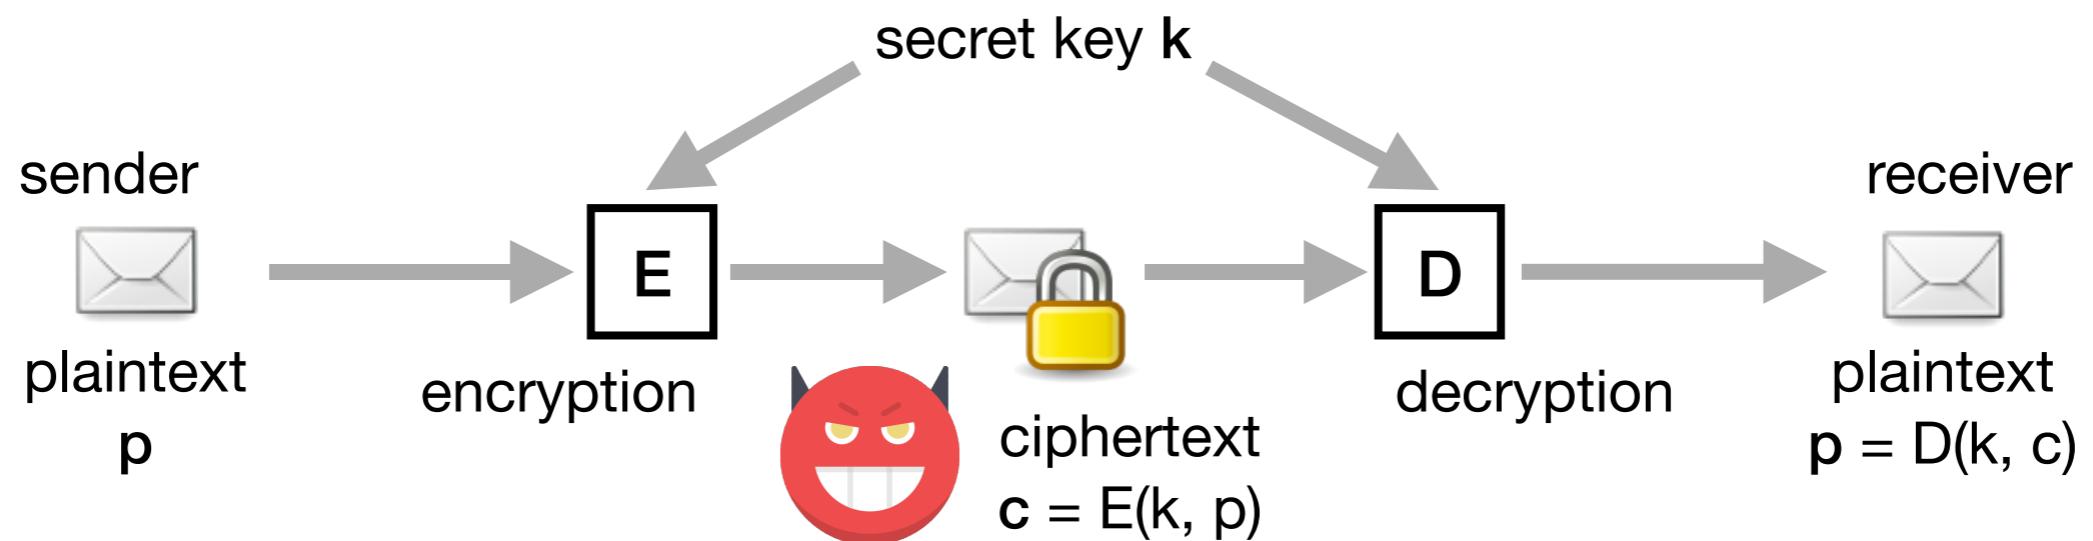
How to generate truly random keys on a computer?

- Clock drift
 - modern computers often have more than one crystal oscillator for timing (e.g., for real-time clock, for CPU)
 - clock crystals are not precise, and there can be **random variations in their speeds**
- Disk drives
 - hard disk drives have small **random fluctuations in their rotational speeds** due to chaotic air turbulence
 - measuring disk seek-time captures this randomness
- Linux: `/dev/random`
 - multiple sources (e.g., mouse and keyboard activity, disk I/O operations, specific interrupts)
 - similar sources on other operating systems
- ...

Reminder:

Encryption

- Model
 - plaintext **p**
 - secret key **k**
 - encryption algorithm (cipher) **E**
 - ciphertext **c** = $E(k, p)$
 - decryption algorithm (cipher) **D**



- Attacker's goal: recover the secret key or the plaintext
- None of the classic cryptographic algorithms are secure

How to define security?

How to define security?

- First idea:
“Attacker cannot recover the secret key”
 - $E(k, p) = p$ would be considered secure under this definition
- Second idea:
“Attacker cannot recover the complete plaintext”
 - suppose that $E(k, p)$ is secure
 - then, $E'(k, p_1 | p_2) = p_1 | E(k, p_2)$ would also be secure
- Third idea:
“Attacker cannot recover any part of the plaintext”
 - what if some part of the plaintext is deterministic (e.g., “GET / HTTP/1.1 ...”)?

None of these ideas define security correctly

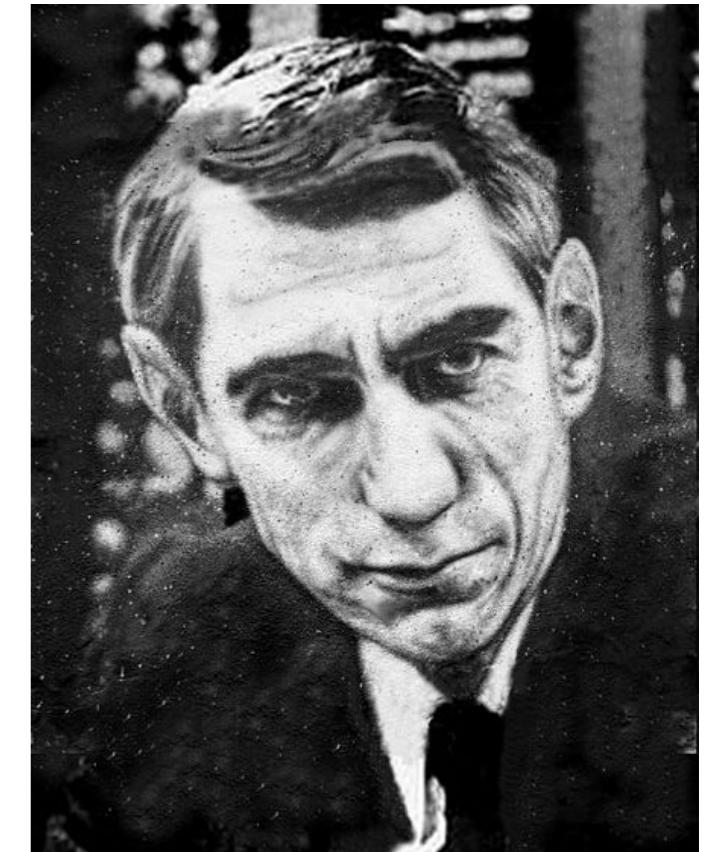


Perfect Security

- Published by Claude Shannon in 1949, based on his classified report from 1945

- **Perfect security:**

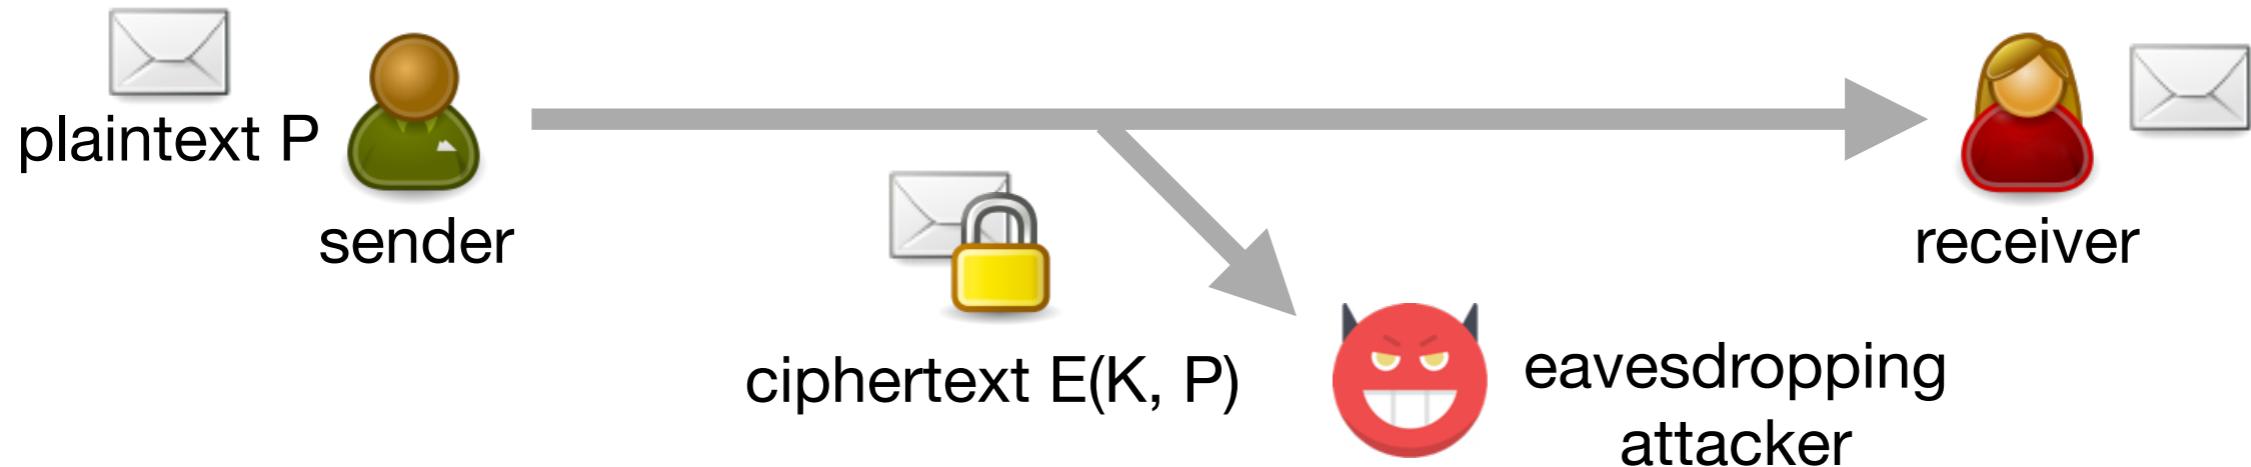
“after a cryptogram is intercepted by the enemy, the a posteriori probabilities of this cryptogram representing various messages be identically the same as the a priori probabilities of the same messages before the interception”



Claude Shannon

≈ attacker gains no information about the plaintext from observing the ciphertext

Perfect Security Formally



- Perfect security

$$\Pr[P = p] = \Pr[P = p | E(K, P) = c]$$

*probability that plaintext P
is a particular message p
(without observing the ciphertext)*

*probability that plaintext P
is a particular message p
given that the ciphertext E(K, P) is c*

- Equivalent to saying that
 plaintext P and ciphertext E(K, P) are independent

One-Time Pad

- Works for binary data as well as alphabetic text
 - plaintext: sequence of bits or letters
- Key: sequence of bits or letters
 1. at least as long as the plaintext
 2. chosen uniformly at random
 3. used to encrypt only one message
- Encryption / decryption: modulo add / subtract each bit (or letter) of the key to the corresponding bit (or letter) of the plaintext / ciphertext
 - binary: $c_i = p_i \oplus k_i$ (XOR operation, i.e., modulo 2 addition)
 - alphabetic: $c_i = p_i + k_i \text{ mod } 26$ and $p_i = c_i - k_i \text{ mod } 26$ (for 26 letters)



One-Time Pad

Binary Example

- XOR operation:

| X | Y | $X \oplus Y$ |
|---|---|--------------|
| 0 | 0 | 0 |
| 0 | 1 | 1 |
| 1 | 0 | 1 |
| 1 | 1 | 0 |

- same as $x + y \bmod 2$ addition (or subtraction since $-1 = 1 \bmod 2$)
- Encryption
 - *plaintext:* 0 1 0 1
 - *key:* 0 1 1 0
 - *ciphertext:* 0 0 1 1
- Decryption
 - *ciphertext:* 0 0 1 1
 - *key:* 0 1 1 0
 - *plaintext:* 0 1 0 1

One-Time Pad

Alphabet Example

- Encode space as the 27th letter and perform addition modulo 27

| | | | | | | | | | | | | | | | | | | | | | | | | | | |
|---|---|---|---|---|---|---|---|---|---|----|----|----|----|----|----|----|----|----|----|----|----|----|----|----|----|--------------|
| a | b | c | d | e | f | g | h | i | j | k | l | m | n | o | p | q | r | s | t | u | v | w | x | y | z | SPACE |
| 0 | 1 | 2 | 3 | 4 | 5 | 6 | 7 | 8 | 9 | 10 | 11 | 12 | 13 | 14 | 15 | 16 | 17 | 18 | 19 | 20 | 21 | 22 | 23 | 24 | 25 | 26 |

- plaintext: mr mustard with the candlestick in the hall
key: pxlmvmsydoфuyrvzwc tnlebnecvgdupahfzzlmnyih
ciphertext: ANKYODKYUREPFJBYOJDSPLREYIUNOFDOIUFPLUYTS
- plaintext: miss scarlet with the knife in the library
key: pftgpmiydgaxgoufhkllmhsqdqogtewbqfgyovuhwt
ciphertext: ANKYODKYUREPFJBYOJDSPLREYIUNOFDOIUFPLUYTS
- Both keys are equally likely
→ attacker cannot learn which key and plaintext are the correct ones from observing the ciphertext



Security of One-Time Pad

- One-time pad is **perfectly secure**
- Proving perfect security:
we have to show that for any plaintext p and ciphertext c ,
$$\Pr[P = p | E(K, P) = c] = \Pr[P = p]$$
- *Proof for special case:* plaintext is 1-bit long (either 0 or 1)
 - key and ciphertext are also 1-bit long
 - given that the **ciphertext** is 0, the probability that the **plaintext** is 0:
$$\Pr[P = 0 | E(K, 0) = 0] = \Pr[P = 0 | K \oplus 0 = 0] = \Pr[P = 0 | K = 0] = \Pr[P = 0]$$
 - given that the **ciphertext** is 1, the probability that the **plaintext** is 1:
$$\Pr[P = 1 | E(K, 1) = 0] = \Pr[P = 1 | K \oplus 1 = 0] = \Pr[P = 1 | K = 1] = \Pr[P = 1]$$
 - ...
 - similar argument works for the general case

One-Time Pad in Practice

- One-time pad has been used in special applications since early 20th century
 - early systems used pencil, paper, and mental arithmetic
 - reportedly used by various intelligence agencies and diplomatic services
- During World War II and the Cold War, Soviet security agencies made heavy use of one-time pads
 - keys were often printed on miniaturized pads made of highly-flammable paper
- Washington-Moscow hotline
 - established in 1963 between the Pentagon and the Kremlin



- each country delivered keys on tapes via its embassy abroad

Limitations of Perfect Security

Why is one-time pad not the “end of cryptography?”

- For perfect security, we always need that
 $\text{length of key} \geq \text{length of plaintext}$
- *Practical problems*
 - distributing and storing long keys
 - generating long truly random keys
- Practical protocols, such as SSL (HTTPS), IPSec, or SSH,
do not use one-time pad

Perfect security
(undefeatable)



“Perfect attacker”
(unbounded capabilities,
except for knowing the key)

Practical security

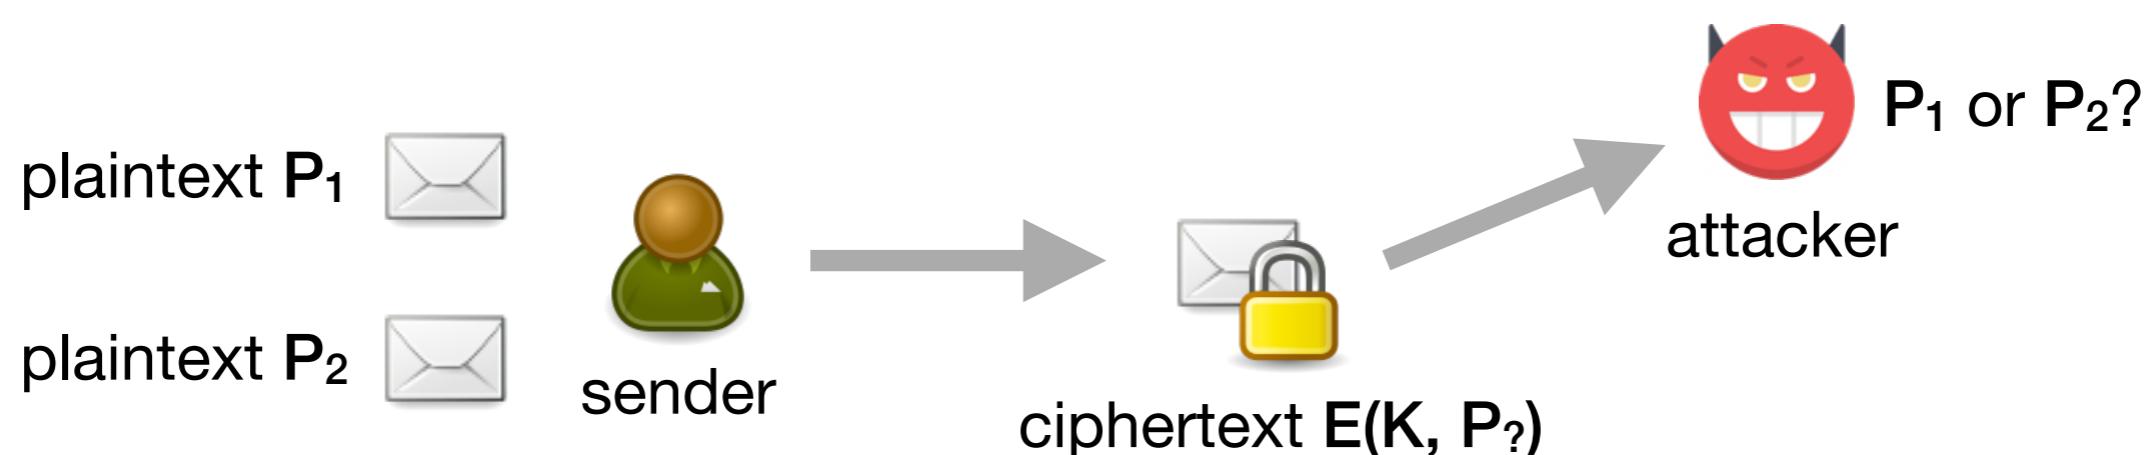


Practical attacker
(bounded computational power)

Semantic Security

- Proposed by Goldwasser and Micali in 1982
- Definition for chosen plaintext attacks (simplified):
 1. attacker chooses two plaintexts
 2. sender encrypts one of the two plaintexts (chosen at random)
 3. attacker observes ciphertext and must guess which plaintext was encrypted

encryption is **semantically secure** if the attacker's advantage for any **efficiently computable** guess is **negligible** over random guessing



Bad Idea:

“Many-Time Pad”

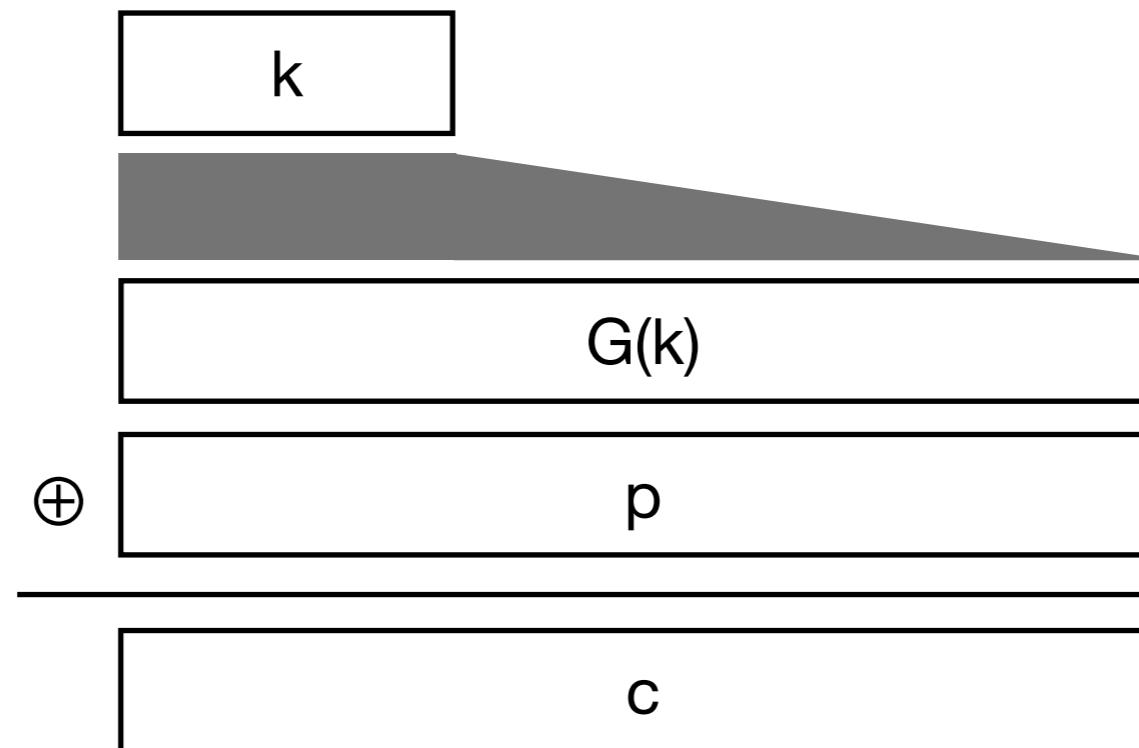
- *Idea:* simply reuse the key for multiple plaintexts
- *Example*
 - key: $k = 01010011$ $k = 01010011$
 - plaintexts: $p_1 = 00111101$ $p_2 = 10110000$
 - ciphertexts: $c_1 = 01101110$ $c_2 = 11100011$
 - attacker can get:
 $c_1 \oplus c_2 = 10001101$
 $c_1 \oplus c_2 = (p_1 \oplus k) \oplus (p_2 \oplus k) = p_1 \oplus p_2 \oplus k \oplus k = p_1 \oplus p_2$
- Problem: if the attacker knows some parts of a plaintext, it can recover the same parts of other plaintexts
 - *example:* attacker knows $p_1 = 00111101$ (e.g., standard protocol format)
 - then, attacker can compute $p_1 \oplus c_1 \oplus c_2 = p_1 \oplus p_1 \oplus p_2 = 0 \oplus p_2 = p_2$



Stream Ciphers

Stream Ciphers

- Idea: make one-time pad practical by extending the key securely

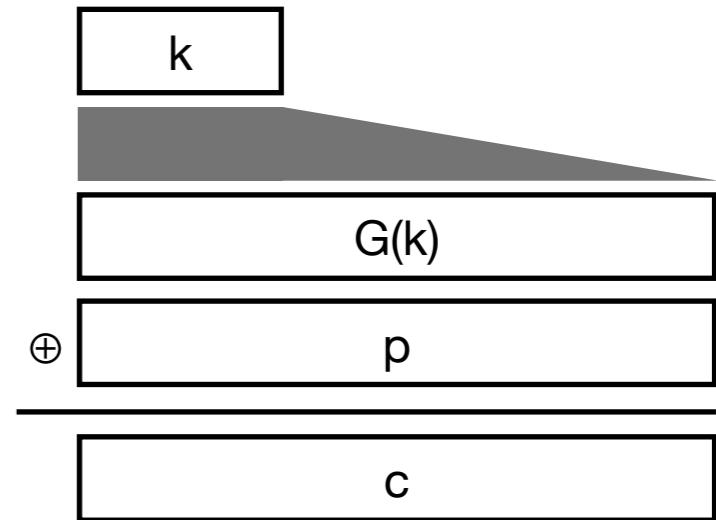


- Deterministic pseudorandom number generator $G(k)$
 - takes a fixed length seed k (i.e., the key) and produces a sequence of bits
- Decryption: generate the same sequence and XOR to the ciphertext

Pseudorandom Number Generator (PRNG)

PRNG: takes fixed length seed, and generates a sequence of bits using a deterministic algorithm

- *Performance requirement*
 - PRNG must generate sequences that are as long as the plaintexts, so it must be **computationally efficient**
 - many algorithms are tailored for hardware architectures or implementations
- *Security requirement*: for an attacker who does not know the key, **bit sequence must be indistinguishable from true randomness**
- Statistical randomness tests (e.g., NIST SP 800-22 “Statistical Test Suite”)
 - frequency of 0s and 1s must be roughly equal
 - distribution of the lengths of uninterrupted “runs” of 0s and 1s (e.g., 10001 or 011110) must be roughly the same as for a true source of randomness
 - ...



passing these tests does not guarantee security!



Bad Example:

Linear Congruential Generator

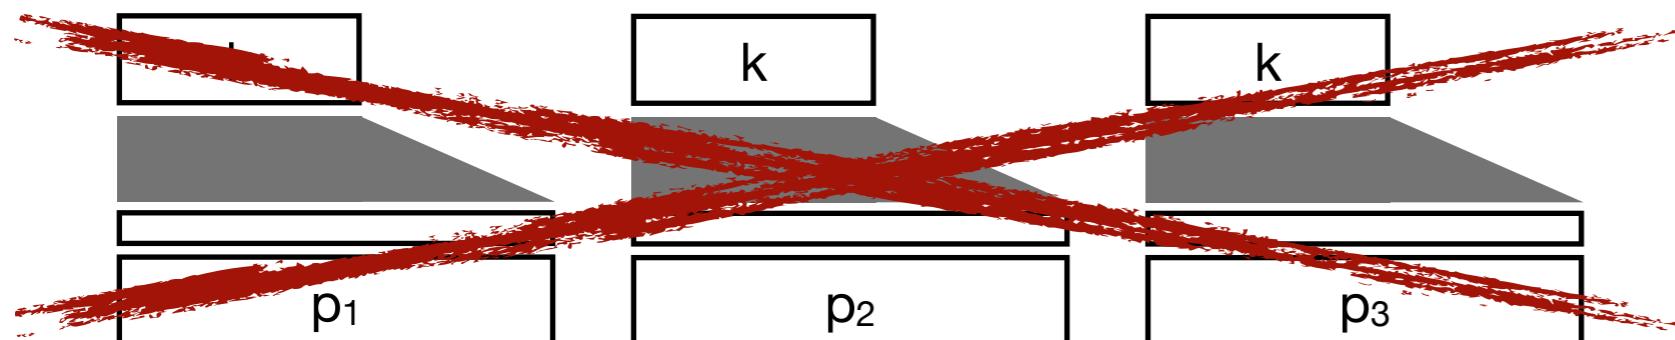
- Generates a sequence of numbers X_1, X_2, \dots from a seed number X_0
$$X_{n+1} = a \cdot X_n + b \pmod{m}$$
- Parameters a , b , and m must be carefully chosen
 - $a = b = 1 \rightarrow$ sequence $0, 1, 2, 3, \dots$
 - $a = 7, b = 0, m = 32$, and $X_0 = 7 \rightarrow$ sequence $7, 17, 23, 1, 7, \dots$
 - if m is prime and $b = 0$ (and a is properly chosen), then **sequence length is $m - 1$**
 - *example:* $X_{n+1} = 7^5 \cdot X_n \pmod{2^{31} - 1}$
- Passes many statistical tests, but provides **no security**
 - if the attacker knows the parameters, it can predict the sequence from one observed value
 - parameters can be efficiently computed from three observed values

Designing Secure Stream Ciphers

- Cryptanalytic attacks
 - attacker may know some bits of the plaintext → may know some bits of the sequence
 - **goal:** bit sequence must be indistinguishable from true randomness
 - security requirement
 1. **uniform distribution:**
frequency of 0s and 1s is approximately equal in the generated sequence
 2. **independence:**
no subsequence can be inferred from another, disjoint subsequence
- Brute-force attacks
 - N-bit long key can take 2^N different values → attacker may try all of them
 - since 2014, the National Institute of Standard and Technology (NIST) recommends at least **112-bit** keys for encryption
 - as computers get faster over time, key sizes must be increased
- Key re-use: *can we use the same key to encrypt multiple plaintexts?*

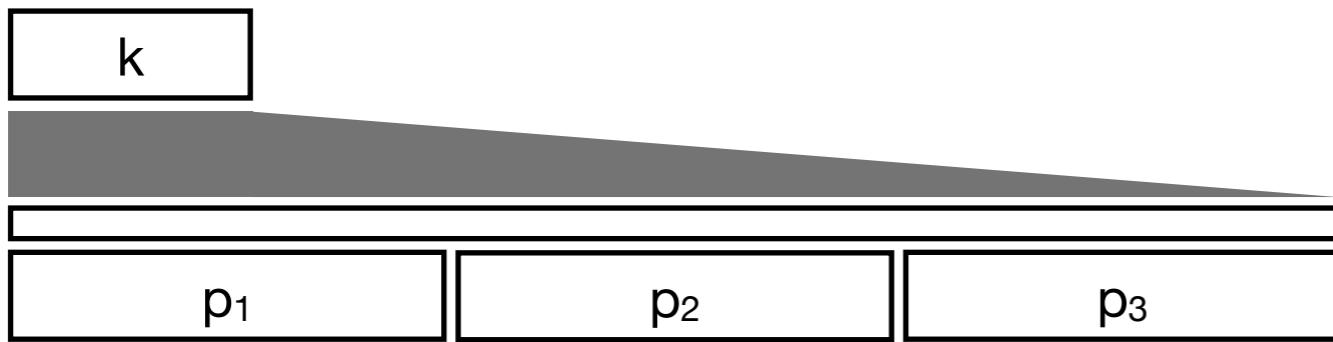
Key Reuse

- Naive approach ~ “many-time pad”



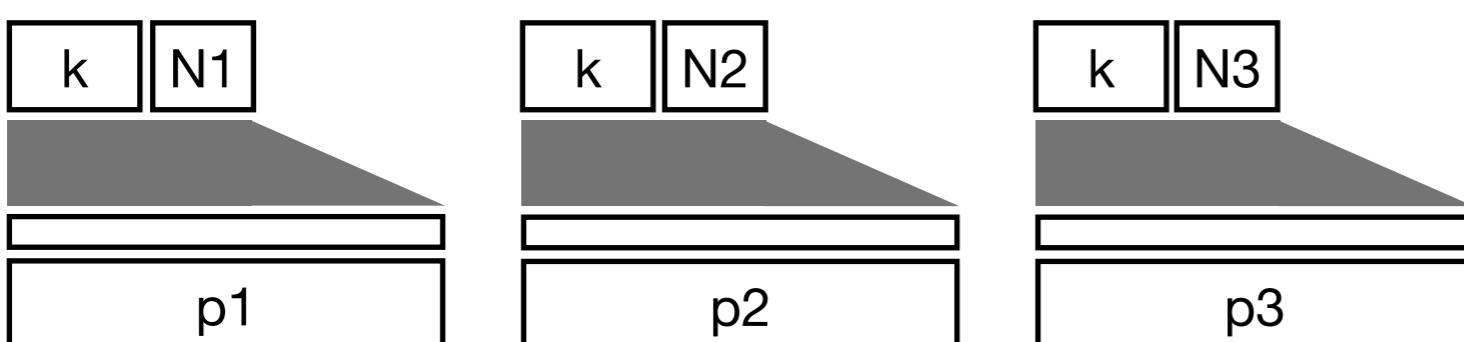
- attacker learns $p_1 \oplus p_2$, $p_2 \oplus p_3$, $p_1 \oplus p_3$, etc.
- if the attacker knows one plaintext, it can recover the others

- One continuous sequence for all plaintexts



- stream cipher must support seeking to any position in the sequence
- otherwise, full sequence must be generated from beginning to decrypt p_3

- Nonces



- nonce: number used once
- mix (e.g., prefix or XOR) key with each nonce → effectively different key for each plaintext

Next lecture:

Block Ciphers